



ENCS 6921 - Industrial Stage and Training

Winter 2023

April Monthly Report

Final Report

Submitted By:

Manan Dineshkumar Paruthi - 40192620

Submitted To:

Dr. Rajagopalan Jayakumar

**Director, Co-op Program, Associate Professor, Computer Science and Software
Engineering**

Introduction & Problem Statement

Before starting my Winter 2023 Coop, I had a discussion with my manager about the requirements of this course i.e ENCS 6921 Industrial Stage and Training and what can be done to aim to work above and beyond the scope of my internship.

So, he informed me that application logging is very important and critical in the industry, for Identifying and debugging issues in the production environment. As it contains sensitive and personal customer data, live debugging can not be done. Also, analysis of logs data helps us to gain information which is important to improve the quality of the applications.

Currently, the company is using Splunk and they are willing to migrate it to ELK Stack (Elastic Search, Logstash & Kibana) as it is open source so free of cost while Splunk is costly and requires a license.

But the issue with ELK Stack is that it is more difficult to set up and maintain. So, it would be great if I can make a POC (Proof Of Concepts) on the same which can be demoed to the senior management and based on that we can start the migration to ELK Stack.

The main reasons for migration are :

- 1) Cost: ELK Stack is open-source, which means it is free to use, whereas Splunk is proprietary software and requires a license, it makes ELK a more cost-effective option for organisations and helps them to reduce the budget expenditure in this economic crisis and layoff environment.
- 2) Flexibility: ELK Stack is highly customizable and can be tailored to specific needs, whereas Splunk's capabilities are more limited. With ELK, organizations can choose to use different components and add-ons, and can also create custom plugins to extend their functionality
- 3) Data Ownership: With ELK Stack, organizations have more control over their data and can decide where it is stored and how it is processed, whereas Splunk stores data on its own servers.

What is ELK Stack ?

The ELK Stack, consisting of Elasticsearch, Logstash, and Kibana, is a popular open-source solution for application logging and data analysis. The stack is designed to collect, process, and analyze log data, providing organizations with valuable insights into their systems and applications.

What problems are solved by ELK Stack ?

One of the main problems that the ELK Stack addresses is the ability to search, analyze, and visualize large volumes of log data. Elasticsearch, the core component of the stack, is a powerful search and analytics engine that allows for fast and efficient data querying. Logstash, another component, is a data processing pipeline that can normalize, enrich, and process log data before it is indexed in Elasticsearch. Kibana, the final component, is a visualization tool that allows users to create interactive dashboards and charts to analyze the data.

What are the advantages of ELK Stack ?

The ELK Stack has several advantages that make it a popular choice for application logging and analysis. One of the main advantages is its open-source nature, which allows for greater flexibility and customization. Additionally, the stack has a large and active community, which provides a wealth of resources and support. The stack is also highly scalable and can handle large volumes of data, which makes it well-suited for organizations with large and complex systems.

What are the disadvantages of ELK Stack ?

However, the ELK Stack also has some drawbacks. One of the main cons is that it can be more difficult to set up and maintain compared to other solutions. Additionally, the stack does not have a built-in alerting system, which can make it difficult to detect and respond to issues in a timely manner.

What is the demand for ELK Stack in the industry ?

The demand for the ELK Stack in the industry is high, with many organizations using the stack for application logging and analysis. The stack is well suited for organizations that have large and complex systems and that need more control over their data. Additionally, the stack is popular among organizations that are looking for an open-source solution that can be customized to their specific needs.

Comparison of ELK Stack with Splunk

When compared to other solutions, the ELK Stack is often compared to Splunk, which is a proprietary log management and data analysis tool. While both solutions offer similar capabilities, the ELK Stack is open-source and offers greater flexibility and customization. However, Splunk is known for its ease of use and intuitive interface and offers a wide range of pre-built visualizations and alerts. Both solutions have their own strengths and weaknesses, and the choice between them will depend on the specific needs and resources of the organization.

Project Scope

The MEAN Stack application (MongoDB, ExpressJS, Angular, NodeJS), a full stack application with MongoDB as NoSQL database, Angular as Frontend and ExpressJS & NodeJS as Backend, which will have create account, login & logout features along with logging functionality. And logs parsing and analysis functionality by using ELK stack.

It will simulate an industry-level application environment which includes an application along with logging, parsing, and analysis capabilities using open-source technologies.

Solution Approach

This project is a web application that focuses on log management and data analysis. The application allows users to create accounts, log in to their accounts, and log out of their accounts. For every user action, events are logged, and log data is generated. This log data is then parsed and used to analyze and visualize trends and metrics using an informative dashboard.

Log management technologies enable businesses to extract important information from unstructured data such as response times, customer plan dollar values, and performance resource usage information. These technologies roll this information up into metrics dashboards that give businesses insights into a wide range of trends that are happening across their systems and operations.

One of the advantages of log data is that it always maintains evidence of what happened. For example, if there is a sudden spike in the number of signups, businesses can quickly validate the change by checking the logs to see who signed up and when. Traditional metrics dashboards do not always maintain the source of data used to calculate the metrics, so validating sudden changes often requires a discussion with engineers, which can take days to complete.

By using this web application, businesses can gain insights into their operations by analyzing log data. The dashboard provides a visual representation of the data, making it easier for businesses to identify trends and patterns. The application makes it easy to log user actions and analyze the data, providing valuable insights for businesses to improve their operations and performance.

In summary, this project focuses on log management and data analysis, allowing businesses to gain insights from their log data. The application provides a user-friendly dashboard for visualizing the data, making it easier for businesses to identify trends and patterns. By using this application, businesses can make data-driven decisions to improve their operations and performance.

It is a web application that utilizes the MEAN stack for front-end and back-end development and the ELK stack for log management and data analysis.


The MEAN stack consists of four technologies: MongoDB, ExpressJS, AngularJS, and NodeJS. MongoDB is a NoSQL document-oriented database used to store and manage data. ExpressJS is a web application framework used to build APIs and web applications. AngularJS is a client-side framework used to build dynamic web applications. NodeJS is a server-side JavaScript runtime used to build scalable and efficient server-side applications.

The front end of the application is built using Angular, a client-side framework that provides a seamless user experience. Angular makes it easy to build dynamic and interactive web applications. The back end of the application is built using NodeJS, ExpressJS, and MongoDB. NodeJS provides a scalable and efficient runtime environment for server-side JavaScript applications, while ExpressJS provides a robust framework for building APIs and web

applications. MongoDB is used to store and manage data in a document-oriented format, making it easy to scale and maintain.

The Winston Logging tool is used for logging and saving application events. Winston is a JavaScript logging library that allows developers to log messages to various destinations such as the console, a file, or a database. By logging events, developers can analyze and assess threats and analyze errors before they disrupt broader business workflows.

Cypress is a front-end testing tool used to perform different actions that generate logs. These logs can be used for data analysis and visualization. Cypress is a JavaScript-based testing framework that allows developers to write and run automated tests for web applications.



The ELK stack is used to collect, store, and analyze log data from the application. The ELK stack consists of three technologies: Elasticsearch, Logstash, and Kibana. Logstash is used to collect log data from various sources, Elasticsearch is used to store the log data, and Kibana is used to provide visualizations and analytics. The ELK stack makes it easy to collect and analyze log data, providing insights into the performance and behaviour of the application.

Technologies Used



1) MongoDB

MongoDB is a popular NoSQL document-oriented database that stores data in a flexible, JSON-like format called BSON (Binary JSON). It is an open-source database that provides high scalability and performance for modern web and mobile applications.

Unlike traditional relational databases that use tables to store data, MongoDB stores data in collections and documents. Documents are like rows in a table, but with a much more flexible schema. This means that data can be added, removed, or modified without altering the existing structure of the database.

MongoDB is known for its high scalability and performance, making it ideal for handling large amounts of unstructured data. It provides a built-in sharding and replication mechanism that ensures high availability and scalability. Sharding distributes data across multiple servers, allowing applications to handle more data and traffic. Replication ensures that data is replicated to multiple servers, providing automatic failover and recovery in the event of a server failure.

In addition, MongoDB has a powerful query language that supports complex queries and aggregation operations. It also supports full-text search, geospatial indexing, and graph processing.

MongoDB is used by many organizations across different industries, from startups to large enterprises, for a variety of use cases such as e-commerce, social networking, content management, and IoT. It provides developers with a flexible and scalable database solution that can adapt to their changing needs.

Advantages of MongoDB:

1. **Scalability:** MongoDB can easily scale horizontally by adding more servers to the cluster. This makes it a great choice for applications with rapidly growing data volumes and high traffic loads.

2. **Flexible Data Model:** MongoDB allows for the storage of unstructured data, making it ideal for applications that have variable data requirements. The JSON-like data model allows for easy integration with web applications and provides a more natural way to store and retrieve data.
3. **High Performance:** MongoDB provides high-performance data retrieval and query processing, with support for a variety of indexing options, including geospatial and text indexes.
4. **Replication and High Availability:** MongoDB provides automatic replication and failover capabilities, ensuring high availability and data durability.
5. **Open Source:** MongoDB is open-source software, which means it is free to use and comes with a large community of developers and users who contribute to its development and support.

Disadvantages of MongoDB:

1. **Data Consistency:** MongoDB sacrifices data consistency for scalability and performance. This means that in certain situations, data may not be consistent across all nodes in the cluster.
2. **Limited Transaction Support:** MongoDB does not support transactions across multiple documents, which can make it challenging to ensure data consistency across a complex system.
3. **Limited SQL Support:** MongoDB does not support SQL, which may make it more challenging for users who are accustomed to working with relational databases.

Use Case Scenarios:

1. **E-commerce Applications:** MongoDB is ideal for e-commerce applications that need to store and manage large volumes of product data, including product descriptions, images, and pricing information.
2. **Social Media Applications:** MongoDB can be used to store and manage user data for social media applications, including user profiles, posts, and comments.
3. **Internet of Things (IoT):** MongoDB is well-suited for IoT applications that collect and process large volumes of sensor data, including temperature readings, GPS locations, and other sensor data.
4. **Content Management Systems (CMS):** MongoDB can be used to store and manage content for CMS applications, including blog posts, articles, and multimedia files.
5. **Real-time Analytics:** MongoDB can be used for real-time analytics applications, including log analysis and machine data processing, where the ability to scale horizontally and handle large volumes of unstructured data is critical.

2) Express JS

ExpressJS is a popular open-source web application framework for Node.js. It provides a robust set of features for creating web applications and APIs, including middleware, routing, and templating engines.

With its minimalist approach to web development, ExpressJS allows developers to build scalable and flexible web applications with ease. It provides a lightweight and unopinionated structure, allowing developers to choose their preferred tools and libraries.

One of the key features of ExpressJS is its middleware system, which allows developers to add additional functionality to the application. Middleware can be used for handling requests, performing authentication, logging, and more.

ExpressJS also provides a routing mechanism that enables developers to define routes for their applications. These routes can be used to handle different HTTP requests, such as GET, POST, PUT, and DELETE, and map them to specific functions or handlers.

In addition, ExpressJS supports a variety of templating engines, such as EJS, Pug, and Handlebars, which can be used to generate dynamic HTML content. It also provides support for static file serving, which can be useful for serving images, CSS, and JavaScript files.

Overall, ExpressJS is a powerful and flexible web framework that is widely used by developers around the world. Its simplicity, scalability, and wide range of features make it a popular choice for building web applications and APIs.

Advantages of ExpressJS:

1. **Flexibility:** ExpressJS is a minimalist framework, and it does not impose any particular coding style, making it highly flexible. Developers can use it to build web applications using their preferred libraries and tools.
2. **Easy to Learn:** ExpressJS is relatively easy to learn, especially for developers who are already familiar with Node.js. Its simple syntax and modular structure make it easy to use and understand.
3. **Wide Community Support:** ExpressJS has a large and active community of developers who contribute to its development and offer support through forums, blogs, and other online resources.
4. **Middleware:** ExpressJS provides a robust middleware system that allows developers to add additional functionality to their web application. Middleware can be used to handle authentication, logging, and other functions.
5. **Scalability:** ExpressJS is highly scalable and can handle a large number of requests without any performance issues.

Disadvantages of ExpressJS:

1. **Steep Learning Curve:** Although ExpressJS is relatively easy to learn, it can be challenging for developers who are new to Node.js or web development.
2. **Lack of Built-in Features:** ExpressJS is a minimalist framework, and it does not provide many built-in features. Developers have to use external libraries to add functionality to their applications.
3. **Security:** As with any web application, security is a significant concern when using ExpressJS. Developers need to ensure that their application is secure and free from vulnerabilities.

Use Cases of ExpressJS:

1. **Building Web Applications:** ExpressJS is ideal for building web applications, especially those that require a lightweight and modular structure.
2. **Creating APIs:** ExpressJS is also used for creating APIs that can be used by other applications or services.
3. **Developing Real-time Applications:** ExpressJS is used in real-time applications that require fast and efficient communication between the server and the client, such as chat applications, online gaming, and social networks.
4. **Creating Microservices:** ExpressJS can also be used for creating microservices, which are small, independent services that work together to form a larger application.

3) Angular

Angular is a popular open-source web application framework developed and maintained by Google. It is written in TypeScript, a superset of JavaScript, and is used for building dynamic, single-page web applications (SPAs).

Angular provides developers with a set of tools and features to simplify the development process and improve the performance of their applications. It uses a component-based architecture where each part of an application is represented by a reusable and encapsulated component. These components can communicate with each other through services and observables, making it easier to manage and maintain large-scale applications.

One of the best features of Angular is its two-way data binding, which allows changes made in the model to be reflected immediately in the view, and vice versa. This reduces the need for developers to write additional code to update the UI and makes the development process more efficient.

Angular also provides features such as dependency injection, which allows components to be loosely coupled and promotes code reusability, and the Angular CLI, which provides a command-line interface to generate, build, and test Angular applications.

Overall, Angular is a powerful and flexible framework that provides developers with the tools and features they need to build complex and scalable web applications. However, it does have a steep learning curve and requires a strong understanding of TypeScript and other web development technologies.

Advantages of Angular:

1. **Robust framework:** Angular is a robust framework that offers powerful features, such as two-way data binding, dependency injection, and reusable components, to simplify the development process and improve the performance of applications.
2. **Large community and support:** Angular has a large and active community of developers who contribute to the framework, provide resources and support, and help solve problems.

3. Code reusability: Angular's component-based architecture and dependency injection make it easy to reuse code, reducing development time and increasing efficiency.
4. Powerful CLI: Angular's CLI provides a powerful command-line interface that allows developers to generate, build, and test Angular applications easily and efficiently.
5. Scalability: Angular's architecture and design patterns make it a great choice for building large-scale and complex applications.

Disadvantages of Angular:

1. Steep learning curve: Angular has a steep learning curve and requires a strong understanding of TypeScript and other web development technologies.
2. Performance issues: Angular's large size and complexity can lead to performance issues, especially when building small and simple applications.
3. Limited flexibility: Angular's opinionated approach and strict guidelines can limit flexibility and make it difficult to customize the framework.

Use case scenarios of Angular:

1. Enterprise applications: Angular is a popular choice for building enterprise applications due to its robustness, scalability, and code reusability.
2. Single-page applications: Angular's component-based architecture and powerful features make it an ideal framework for building single-page applications.
3. Real-time applications: Angular's two-way data binding and observables make it a great choice for building real-time applications, such as chat applications and dashboards.
4. E-commerce websites: Angular's scalability, flexibility, and powerful CLI make it an ideal choice for building e-commerce websites that require complex features and high performance.
5. Hybrid mobile applications: Angular can be used with Ionic, a popular mobile app development framework, to build hybrid mobile applications that work across multiple platforms.

4) Node JS

NodeJS is an open-source, cross-platform JavaScript runtime environment that allows developers to build server-side applications using JavaScript. It was developed by Ryan Dahl in 2009 and has since grown in popularity among developers due to its ease of use, speed, and versatility.

One of the key advantages of NodeJS is that it is based on the V8 engine, which is the same engine used by Google Chrome. This means that NodeJS can execute JavaScript code very quickly, making it ideal for building high-performance applications. Additionally, NodeJS uses an event-driven, non-blocking I/O model, which allows it to handle large volumes of requests without blocking other processes.

NodeJS comes with a rich library of modules and packages, making it easy to extend its capabilities and build complex applications. These modules cover a wide range of functionality, including HTTP and HTTPS servers, file system access, database connectivity, and much more.

Additionally, the npm (Node Package Manager) repository contains over a million packages that developers can use to enhance their applications.

NodeJS is often used for building real-time applications such as chat applications, online gaming platforms, and collaborative tools. It is also popular for building RESTful APIs and serverless applications. NodeJS can be run on a variety of platforms, including Windows, Mac OS, and Linux, making it a versatile option for developers.

In summary, NodeJS is a powerful and versatile tool that allows developers to build fast, scalable server-side applications using JavaScript. Its ease of use, speed, and large library of modules make it a popular choice for a wide range of applications.

Advantages of NodeJS:

1. **Fast and scalable:** NodeJS is built on an event-driven, non-blocking I/O model that makes it fast and efficient, allowing it to handle a large volume of requests with ease.
2. **Cross-platform compatibility:** NodeJS can be run on multiple platforms, including Windows, Mac OS, and Linux, making it a versatile option for developers.
3. **Large community and ecosystem:** NodeJS has a vast community of developers and a rich library of modules and packages available through the npm repository, making it easy to extend its capabilities.
4. **JavaScript-based:** NodeJS uses JavaScript, which is a widely used language, making it easy for developers to learn and use it.
5. **Ideal for real-time applications:** NodeJS is great for building real-time applications such as chat applications, online gaming platforms, and collaborative tools.

Disadvantages of NodeJS:

1. **Single-threaded:** NodeJS is single-threaded, which means it can only handle one request at a time, making it unsuitable for CPU-intensive tasks.
2. **Not suitable for heavy computation:** NodeJS is not well suited for performing heavy computational tasks such as machine learning or scientific computing.
3. **Callbacks can be challenging:** NodeJS relies heavily on callbacks, which can make it difficult to write and maintain code.
4. **Not ideal for database-intensive applications:** NodeJS may not be the best choice for database-intensive applications as it may have to deal with multiple I/O operations, which could result in slower performance.
5. **Immature tooling:** As NodeJS is relatively new, some of the toolings around it may not be as mature as other languages or frameworks.

Use case scenarios of NodeJS:

1. **Real-time applications:** NodeJS is ideal for building real-time applications such as chat applications, online gaming platforms, and collaborative tools.
2. **RESTful APIs:** NodeJS is great for building RESTful APIs as it can handle a large number of requests efficiently.

3. Microservices: NodeJS can be used to build microservices as it can easily be integrated with other services.
4. Single-page applications: NodeJS is suitable for building single-page applications as it can render content quickly and efficiently.
5. Streaming services: NodeJS is ideal for building streaming services as it can handle large volumes of data in real-time.

5) Cypress

Cypress is an open-source test automation tool used for web applications. It is a JavaScript-based testing framework that allows developers to write automated tests for the front-end of web applications.

Cypress provides a comprehensive suite of testing capabilities, including automated testing, end-to-end testing, integration testing, and unit testing. It uses a unique architecture that runs the test code in the same runtime as the application being tested, allowing for real-time debugging, automatic reloading, and quicker test feedback.

One of the significant advantages of Cypress is its ability to speed up the testing process. The tool has a fast test execution time and requires minimal setup and configuration, allowing developers to quickly test their code changes and catch bugs before they become more significant issues.

Cypress also has an easy-to-use dashboard that provides a real-time view of the test runs, including the results, logs, and screenshots, making it easy to identify issues and debug them quickly.

Overall, Cypress is a powerful tool that simplifies the testing process, saves time, and increases the efficiency of web application testing. Its user-friendly interface and comprehensive testing capabilities make it an excellent choice for developers looking to streamline their testing workflow.

Advantages of Cypress:

1. Easy to use and learn: Cypress is designed to be easy to use and learn for developers. It has an intuitive interface that allows developers to write tests quickly and efficiently.
2. Fast test execution: Cypress is known for its fast test execution time. It runs tests in the same runtime as the application being tested, which results in quicker feedback for developers.
3. Real-time reloading and debugging: Cypress allows developers to debug their tests in real-time and automatically reload the application being tested. This feature makes it easier to identify and fix issues quickly.
4. Comprehensive testing capabilities: Cypress provides a wide range of testing capabilities, including automated testing, end-to-end testing, integration testing, and unit testing.
5. Built-in assertions and mocking: Cypress comes with built-in assertions and mocking capabilities that make it easier to write tests and simulate different scenarios.

Disadvantages of Cypress:

1. Limited browser support: Cypress only supports Chrome and Firefox browsers, which can be limiting for some developers.
2. Limited mobile device support: Cypress does not support mobile devices, which can be a significant disadvantage for applications that require testing on mobile devices.
3. Limited support for APIs: Cypress is primarily focused on testing web applications and does not provide extensive support for testing APIs.

Use case scenarios for Cypress:

1. End-to-end testing: Cypress is a great tool for end-to-end testing of web applications. It allows developers to simulate real-world scenarios and test the entire application flow.
2. Automated testing: Cypress is an excellent tool for automated testing of web applications. Its fast test execution time and real-time debugging capabilities make it easy to write and execute automated tests.
3. Integration testing: Cypress provides a robust set of integration testing capabilities that allow developers to test the interaction between different components of the application.
4. Unit testing: Cypress can also be used for unit testing of web applications. Its built-in assertions and mocking capabilities make it easier to write and execute unit tests.

6) Winston Logging Tool

Winston is a versatile and popular logging library for Node.js applications. It provides a simple and flexible API for creating and managing logs, allowing developers to easily capture and store application events and errors.

With Winston, developers can create and customize multiple loggers, each with its own transport mechanism for outputting logs to various destinations such as the console, files, databases, or remote servers. Winston also supports a variety of logging levels, from simple information messages to warnings and errors, allowing developers to fine-tune the verbosity of their logs.

Winston also provides a range of built-in features, such as support for logging metadata, customizable message formatting, and support for logging unhandled exceptions. Additionally, Winston has a robust ecosystem of plugins and extensions that can add even more functionality, such as logging to cloud services or integrating with other libraries and frameworks.

Overall, Winston is a powerful and flexible logging library that can be customized to meet the needs of a wide range of Node.js applications. Its rich feature set, ease of use, and active community make it a popular choice for logging in Node.js projects.

Advantages of Winston Logging Tool:

1. **Versatile:** Winston supports multiple transports, including console, file, database, and cloud services, making it a flexible and adaptable logging solution for a variety of use cases.
2. **Customizable:** Winston allows developers to configure logging levels, formatting, and metadata to meet their specific requirements.
3. **Robust:** Winston includes features such as exception handling and log rotation, which make it a reliable and resilient logging solution.
4. **Community Support:** Winston has a large and active community of users and contributors, providing access to plugins and extensions that can add additional functionality and help resolve issues.

Disadvantages of Winston Logging Tool:

1. **Configuration Complexity:** Winston can be complex to configure, especially for complex use cases or for developers who are new to logging.
2. **Performance Overhead:** Depending on the logging configuration, Winston can introduce a performance overhead, which may impact application performance in high traffic scenarios.
3. **Limited Real-Time Analysis:** Winston is primarily designed for capturing logs, and may not be the best solution for real-time analysis or monitoring.

Use Case Scenarios for Winston Logging Tool:

1. **Debugging:** Winston is an effective tool for debugging Node.js applications, allowing developers to capture and analyze logs to identify and resolve issues.
2. **Auditing:** Winston can be used to capture and store application events and errors for auditing and compliance purposes.
3. **Analytics:** Winston can capture custom application metrics, which can be analyzed to understand application performance and usage patterns.
4. **Production Monitoring:** Winston can be used to monitor production environments, alerting developers to issues such as errors or performance degradation.

7) Elastic Search

Elasticsearch is a popular open-source search and analytics engine that is designed to handle large amounts of data in real-time. It is built on top of Apache Lucene, a high-performance search library, and provides a distributed, scalable, and fault-tolerant search infrastructure that can be used for a wide range of use cases, including log analytics, full-text search, and business analytics.

One of the key features of Elasticsearch is its ability to perform full-text search and analysis of structured and unstructured data. It uses a powerful query language that allows users to search for specific terms or phrases, perform complex aggregations and filters, and even perform predictive analytics using machine learning algorithms.

Elasticsearch is also designed to be highly scalable and fault-tolerant, allowing it to handle large amounts of data across multiple nodes in a cluster. It uses a distributed architecture to store and

index data, which enables it to provide fast and reliable search results even as the amount of data grows.

In addition to search and analytics, Elasticsearch also provides a rich set of APIs and plugins that can be used to extend its functionality. It integrates with a wide range of data sources, including databases, messaging systems, and streaming platforms, and provides powerful tools for monitoring and managing the health of the Elasticsearch cluster.

Overall, Elasticsearch is a powerful and flexible search and analytics engine that can be used to solve a wide range of use cases in industries such as e-commerce, finance, healthcare, and more.

Advantages of Elasticsearch:

1. **High Performance:** Elasticsearch is known for its fast search and indexing performance. It uses the Apache Lucene engine, which makes it incredibly fast when querying large datasets.
2. **Distributed Architecture:** Elasticsearch is designed to be distributed, allowing it to scale horizontally across multiple servers or nodes. This makes it easy to scale the system as the amount of data grows.
3. **Full-Text Search:** Elasticsearch has powerful full-text search capabilities that allow users to search for text across multiple fields and documents.
4. **Open-Source:** Elasticsearch is open-source software and is available for free, making it an affordable solution for small businesses and startups.
5. **Integrations:** Elasticsearch integrates with a wide range of tools and platforms, including Hadoop, Logstash, Kibana, and more.

Disadvantages of Elasticsearch:

1. **Complexity:** Elasticsearch can be complex to set up and manage, especially for non-technical users.
2. **Resource-Intensive:** Elasticsearch requires significant resources in terms of storage, memory, and processing power to run effectively, making it a costly option for some organizations.
3. **No Built-in Security:** Elasticsearch does not come with built-in security features and requires additional configuration to secure the system.

Use Case Scenarios for Elasticsearch:

1. **Search and Discovery:** Elasticsearch is commonly used for search and discovery applications, such as e-commerce websites, where users need to find specific products or content quickly.
2. **Log Analysis:** Elasticsearch can be used to analyze logs from servers and applications, providing insights into system performance and troubleshooting issues.
3. **Business Analytics:** Elasticsearch can be used to analyze large datasets for business intelligence and reporting purposes.

4. **Monitoring and Alerting:** Elasticsearch can be used to monitor system metrics and send alerts when certain thresholds are exceeded, allowing organizations to proactively identify and address issues before they become critical.

8) Logstash

Logstash is an open-source, server-side data processing pipeline that allows you to collect, transform, and send data to various destinations. It is a popular tool in the ELK stack, which also includes Elasticsearch and Kibana.

Logstash uses a variety of input plugins to collect data from different sources, such as log files, network traffic, or databases. It then processes and transforms the data using filters, which can manipulate, modify, and structure the data in a variety of ways. Finally, the data is sent to a destination using output plugins, such as Elasticsearch, Kafka, or Amazon S3.

Logstash provides a powerful and flexible way to manage large amounts of data from different sources, making it easier to centralize and analyze your data. It also provides a range of monitoring and troubleshooting tools, allowing you to track the flow of data through your pipeline and identify any errors or issues.

Overall, Logstash is a powerful tool for data processing and management, allowing you to collect, transform, and send data to a variety of destinations with ease.

Advantages of Logstash:

1. **Versatile:** Logstash can collect data from a variety of sources such as logs, metrics, and data streams. It can also be used to transform, enrich and filter the data before sending it to a destination.
2. **Scalable:** Logstash is designed to handle large amounts of data and can be easily scaled horizontally to accommodate the growing data needs of an organization.
3. **Open source:** Logstash is an open-source tool, which means it's freely available and can be customized to meet your specific requirements.
4. **Integration:** Logstash integrates seamlessly with other tools such as Elasticsearch and Kibana, making it an important part of the ELK stack.
5. **Monitoring:** Logstash provides detailed monitoring capabilities, allowing users to monitor the performance of their data pipeline and identify any issues.

Disadvantages of Logstash:

1. **Complexity:** Logstash can be complex to configure and set up, especially for beginners. The configuration files can be lengthy and difficult to understand.
2. **Performance:** Logstash can be resource-intensive, and processing large amounts of data can impact performance. Careful monitoring and resource allocation are necessary to optimize performance.
3. **Dependencies:** Logstash has a number of dependencies, and if these are not properly installed or configured, it can cause issues with data processing.

Use Case Scenarios:

1. **Log Management:** Logstash is commonly used for log management. It can be used to collect and filter log data from different sources and send it to a centralized location for analysis.
2. **Data Pipeline:** Logstash can be used to create data pipelines for processing and enriching data from multiple sources. It can also be used to send data to different destinations such as Elasticsearch, Kafka, or Amazon S3.
3. **Metrics Monitoring:** Logstash can be used to collect and process metrics data from various sources and send it to a monitoring system such as Grafana or Zabbix.
4. **Security Analysis:** Logstash can be used for security analysis, by collecting and processing security event logs from different sources, and forwarding the data to a SIEM system like Elastic Security or Splunk.
5. **IoT Data Processing:** Logstash can be used to process and analyze data from IoT devices, such as sensors and cameras, and send it to a centralized location for further analysis.

9) Kibana

Kibana is an open-source data visualization and analysis tool designed to work with Elasticsearch, a popular distributed search and analytics engine. Kibana provides a user-friendly interface for exploring, visualizing, and analyzing large volumes of data stored in Elasticsearch.

With Kibana, users can create interactive dashboards, graphs, and charts that display real-time data, making it easy to identify trends, patterns, and outliers in large data sets. Users can also use Kibana to search, filter, and analyze data using a wide range of query and aggregation capabilities.

Kibana is often used for log analysis, monitoring, and troubleshooting, as it allows users to quickly identify and resolve issues within large log data sets. It can also be used for business intelligence and data analysis, enabling users to create insightful visualizations and gain deeper insights into their data.

Kibana supports a variety of data sources, including logs, metrics, and structured data, and can be integrated with a wide range of data storage and processing systems. It also offers a range of customization options, allowing users to tailor the interface to their specific needs and preferences.

Overall, Kibana is a powerful and flexible tool that can help users gain valuable insights from their data in a user-friendly and intuitive way.

Advantages of Kibana:

1. **User-friendly Interface:** Kibana provides a simple and intuitive user interface, allowing users to explore, visualize, and analyze data without requiring specialized technical skills.

2. Real-time data analysis: Kibana provides real-time analysis of data, allowing users to monitor and respond to events as they occur.
3. Powerful data visualization: Kibana enables users to create custom dashboards, graphs, and charts, making it easy to identify trends, patterns, and outliers in large data sets.
4. Customizable: Kibana is highly customizable, allowing users to tailor the interface to their specific needs and preferences.
5. Integration with Elasticsearch: Kibana works seamlessly with Elasticsearch, enabling users to leverage the powerful search and analytics capabilities of Elasticsearch.

Disadvantages of Kibana:

1. Requires Elasticsearch: Kibana requires Elasticsearch to function, which may add additional complexity and cost to the deployment.
2. Limited support for non-structured data: Kibana is primarily designed to work with structured data and may have limited support for non-structured data sources.
3. Requires technical expertise: While Kibana provides a user-friendly interface, it still requires some technical expertise to configure and manage.

Use case scenarios for Kibana:

1. Log analysis: Kibana is commonly used for log analysis, allowing users to monitor and analyze logs in real-time to identify and resolve issues.
2. Business Intelligence: Kibana can be used for business intelligence and data analysis, enabling users to create insightful visualizations and gain deeper insights into their data.
3. Cybersecurity: Kibana can be used for cybersecurity analysis, allowing users to monitor and detect security threats in real-time.
4. IT Operations: Kibana can be used for IT operations monitoring, allowing users to monitor system performance and identify issues before they impact end-users.
5. Marketing Analytics: Kibana can be used for marketing analytics, allowing users to analyze and visualize customer data to gain insights into customer behavior and preferences.

Architecture Diagram

Full Stack Application MEAN Stack

Frontend



Features :
Create Account
Log In
Log Out

Backend



APIs :
/createAccount
/login
/logout

Database



Tables :
UserData
UserSessionData

Automation Tool



To create multiple accounts, login and logout in order to generate logs

Log Management & Data Analysis



Search & Analytical Engine for Data



Collect, Parse & Transform Logs

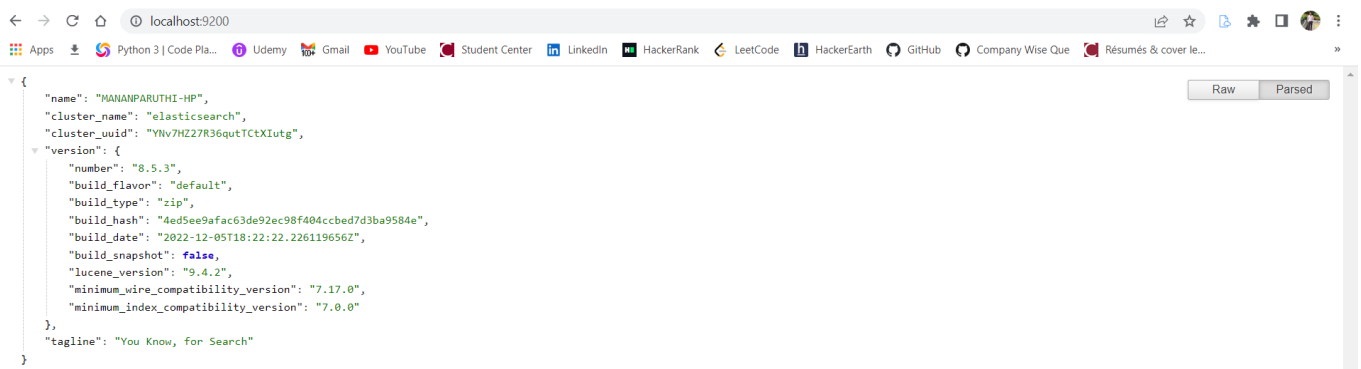


Explore, Visualize & Discover Data

Implementation Details

To set up Elasticsearch on a local system, follow these steps:

- 1) Open the elasticsearch.yml file in the settings folder. To enable automatic index construction, add the following line: "action.auto create index:.monitoring*,.watches,.triggered watches,.watcher-history*,.ml*". Save and exit the file.
- 2) Go to the bin folder and execute "elasticsearch.bat." This will produce a Kibana password and enrollment token. This information should be copied and saved for future use.
- 3) Return to the configuration folder and re-open the elasticsearch.yml file. Under the "xpack.security.http.ssl" and "xpack.security.transport.ssl" sections, change the "enabled" option to false. Save and exit the file.
- 4) Go back to the bin folder and execute "elasticsearch.bat."
- 5) Check "http://localhost:9200/". You will be required to enter your login information. Use the following details: The username is elastic, and the password is the elastic password you copied before. To acquire a new password, navigate to the bin folder and run the command "elasticsearch-reset-password -u elastic."
- 6) Next, navigate to the bin folder and type "elasticsearch-reset-password -u kibana system" to copy and store the Kibana password.

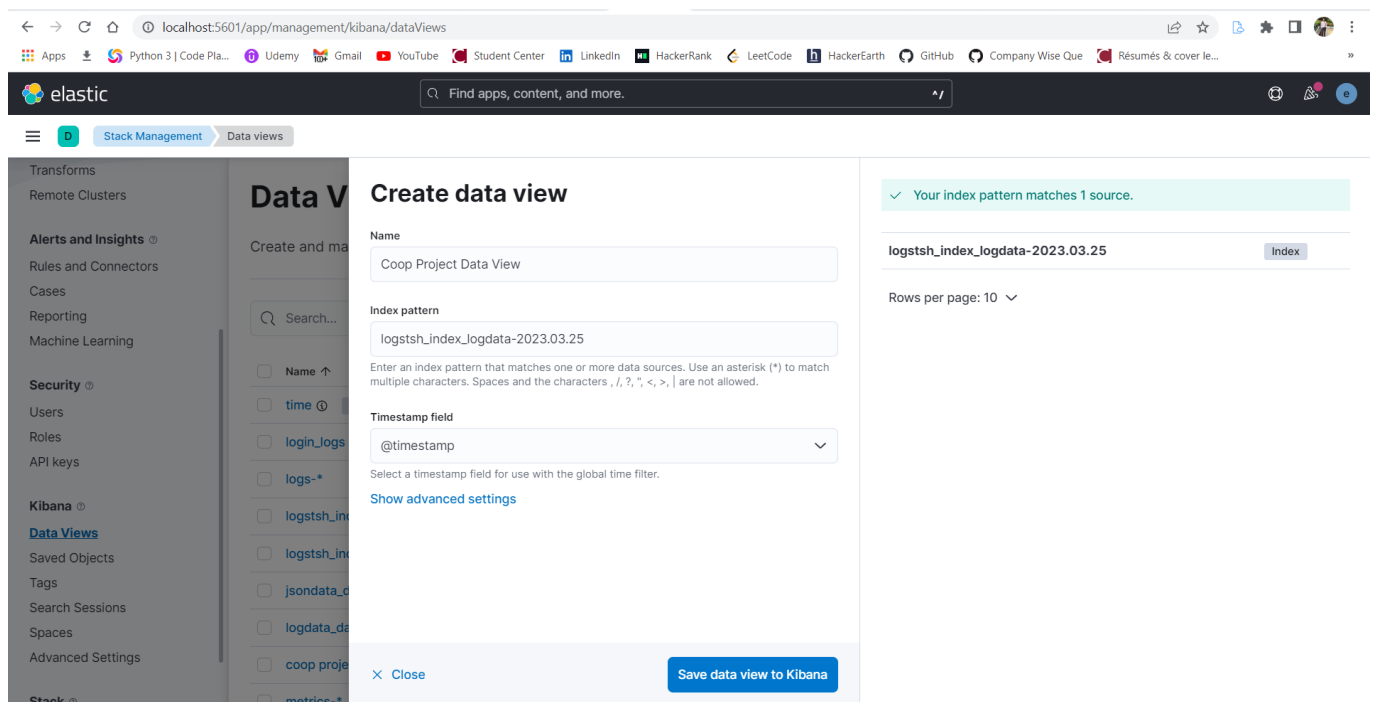


```
{
  "name": "MANANPARUTHI-HP",
  "cluster_name": "elasticsearch",
  "cluster_uuid": "YHv7H27R36qutTctXIutg",
  "version": {
    "number": "8.5.3",
    "build_flavor": "default",
    "build_type": "zip",
    "build_hash": "4ed5ee9afac63de92ec98f404ccbed7d3ba9584e",
    "build_date": "2022-12-05T18:22:22.226119656Z",
    "build_snapshot": false,
    "lucene_version": "9.4.2",
    "minimum_wire_compatibility_version": "7.17.0",
    "minimum_index_compatibility_version": "7.0.0"
  },
  "tagline": "You Know, for Search"
}
```

To set up Kibana on a local system, follow these steps:

- 1) Open the config folder and find the kibana.yml configuration file. Uncomment the server.port, server.host, elasticsearch.host, elasticsearch.username, and elasticsearch.password settings. Replace the password with the Kibana password you already stored. Save and exit the file.
- 2) Go to the bin folder and run the "kibana.bat" programme.

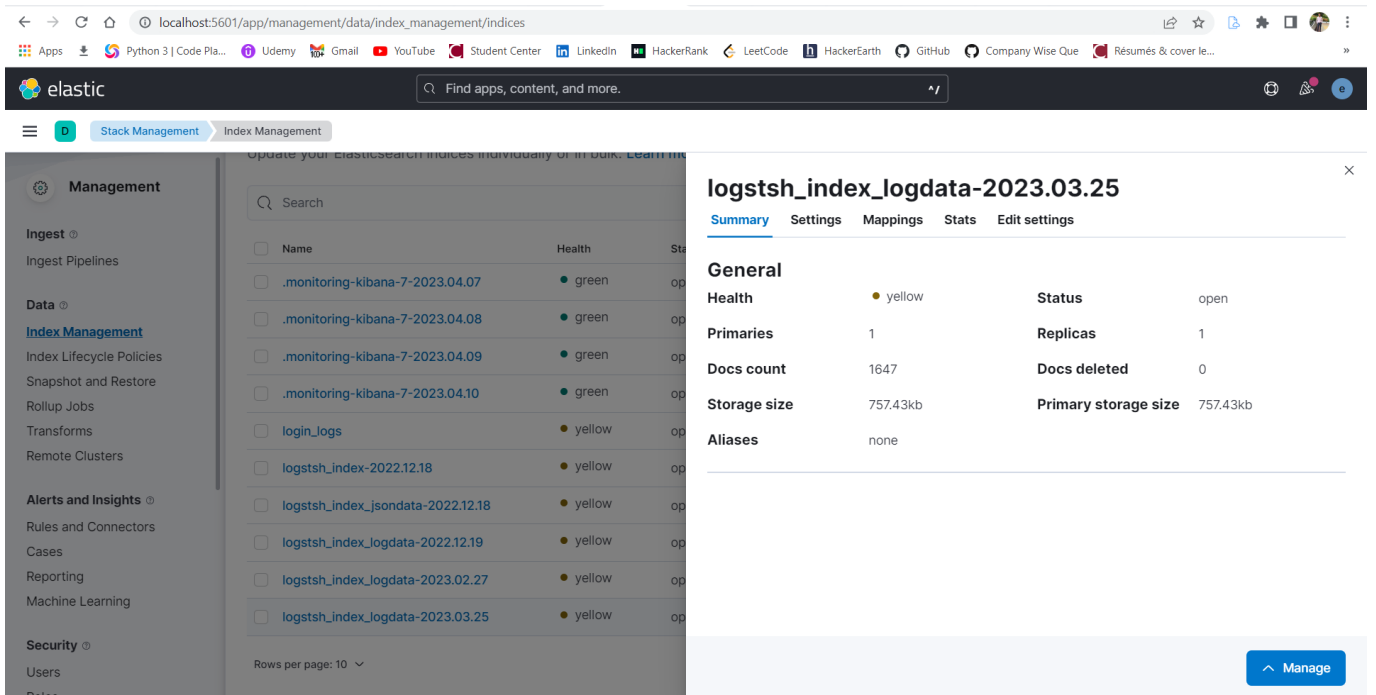
3) Go to "http://localhost:5601." You will be required to enter your login information. Use the following information: The username is "elastic," and the password is the elastic password you previously copied.



To set up Logstash on a local system, follow these steps:

1) Go to the config folder and double-click the learn.conf file to open it. Replace the elastic password with the password you already stored.

2) Go to the bin folder and run the command "logstash -f.config/learn.conf —config.reload.automatic". The supplied configuration file will be used to start Logstash, and the —config.reload.automatic flag will automatically reload the configuration file if changes are detected.



To configure a Kibana dashboard, please follow these steps:

- 1) Go to Stack Management, then Index Management on the Elasticsearch site. Copy the name of the file "logstash index logdata-date>" with today's date.
- 2) After that, navigate to Stack Management and then Data Views. Choose "Create Data View." Enter "logstash index logdata-date>" in the index pattern field and give the data view a name. "Save Data View to Kibana" should be selected.
- 3) Lastly, navigate to the Dashboard area and add a new dashboard. Click "Create Visualization" and choose the data view name you created in the previous step from the selection menu on the upper left. You may then construct data visualisations and upload them to the dashboard.

Grok Filter Implementation

The Grok filter in Logstash provides a means to convert unstructured log data into structured data that can be readily indexed and searched in Elasticsearch. Grok identifies patterns in log data with regular expressions and then assigns those patterns to fields.

The Grok filter extracts fields from several forms of log data, including Apache logs, syslog, and NGINX logs. It may also parse bespoke log formats by specifying a pattern that fits the format.

Grok patterns are set in a configuration file and can be tailored to the log data being processed. The extracted information may then be utilised in Logstash for additional processing or sent directly to Elasticsearch for indexing and searching. The Grok filter is a strong tool for converting unstructured log data into structured data that can be studied and acted upon.

```

filter {
  grok {
    patterns_dir => ["/patterns"]
    # logout, create account, login
    match => { "message" =>
      [
        '%{TIMESTAMP_ISO8601:timestamp} \[%{USERNAME:serviceName}\] %{WORD:logLevel}: %{GREEDYDATA:event} - %{GREEDYDATA:action} : \
{"%{WORD}": "%{WORD:firstName}", "%{WORD}": "%{WORD:lastName}", "%{WORD}": "%{WORD:userName}", "%{WORD}": "%\
{TIMESTAMP_ISO8601:loginDate}", "%{WORD}": "%{NUMBER:sessionTimeInSec}", "%{WORD}": "%{GREEDYDATA:loginCountry}", "%{WORD}": "%\
{WORD:loginBrowser}"\}\}',
        '%{TIMESTAMP_ISO8601:timestamp} \[%{USERNAME:serviceName}\] %{WORD:logLevel}: %{GREEDYDATA:event} - %{GREEDYDATA:action} : \
{"%{WORD}": "%{WORD:firstName}", "%{WORD}": "%{WORD:lastName}", "%{WORD}": "%{WORD:userName}", "%{WORD}": "%{GREEDYDATA:password}"\}\}',
        '%{TIMESTAMP_ISO8601:timestamp} \[%{USERNAME:serviceName}\] %{WORD:logLevel}: %{GREEDYDATA:event} - %{GREEDYDATA:action} : \
{"%{WORD}": "%{WORD:userName}", "%{WORD}": "%{GREEDYDATA:password}"\}\}'
      ]
    }
  }
}

```

Kibana Query Language (KQL) Implementation

Kibana Query Language (KQL) is a simple and easy-to-use language that allows users to use Kibana to search, filter, and analyse Elasticsearch data. KQL is an important component of Kibana since it allows users to explore and visualise data by setting search parameters in a query string.

KQL syntax is built on a field:value format, where the value is a search word and the field is where the search will be done. A query to find all documents that have the phrase "apple" in the "product_name" column, for example, might look like this: "product_name:apple".

KQL has a number of operators that may be used to construct more complicated search queries. Here are a few examples of widely used operators:

1) Wildcard Operator

To find words that meet a pattern, use the wildcard operator (*). A query to search for all documents that have phrases beginning with "app" in the "product_name" column, for example, might look like this: "product_name:app*".

2) Operator of a Range

The range operator may be used to find documents having values that fall inside a certain range. For example, a query to find all documents having "price" values between 10 and 20 might look like this: "price:[10 TO 20]".

3) Boolean Expressions

More complicated search queries may be created by using Boolean operators (AND, OR, NOT). A query to find all documents with the phrase "apple" in the "product_name" column and a value between 10 and 20 in the "price" field, for example, might look like this: "product_name:apple AND price:[10 TO 20]".

4) The Fuzzy Operator

The fuzzy operator () can be used to find texts that include terms that are similar to a certain phrase. A query to search for all documents with phrases comparable to "apple" in the "product_name" column, for example, might look like this: "product_name:apple~".

5) Regular Expressions Operator

To search for documents using regular expressions, use the regular expression operator (/regex/). A query to search for all documents containing phrases that match a regular expression in the "product_name" column, for example, might look like this: "product_name:/^app(le)?/".

KQL also includes a number of functions for doing computations, transformations, and aggregations on data. Here are a few of the most often utilised functions:

1) Mathematical Functions

To execute mathematical operations on data, math functions (abs, ceil, floor, round, and so on) can be utilised. A query to round the values in the "price" column to the closest whole number, for example, might look like this: "round(price)".

2) Date Operations

To execute date-related operations on the data, date functions (date, date_add, date_diff, and so on) can be utilised. A query to find all papers having a timestamp within the previous hour, for example, might look like this: "timestamp:>=now-1h".

3) Functions of Aggregation

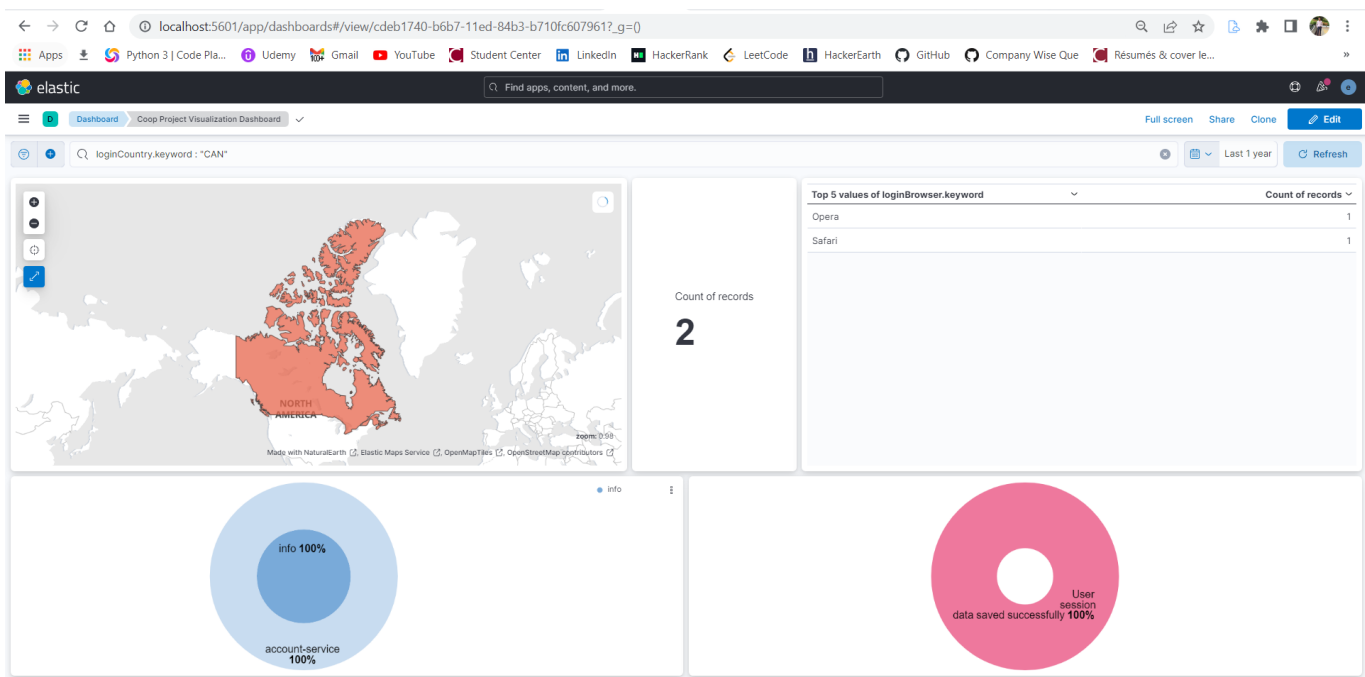
To do statistical computations on data, aggregation functions (count, sum, avg, max, min, and so on) can be employed. A query to determine the average value of the "price" column, for example, might look like this: "avg(price)".

KQL also supports a wide range of data types, including strings, integers, dates, booleans, and arrays. Users can utilise these data types to more clearly describe search criteria and conduct actions on the data.

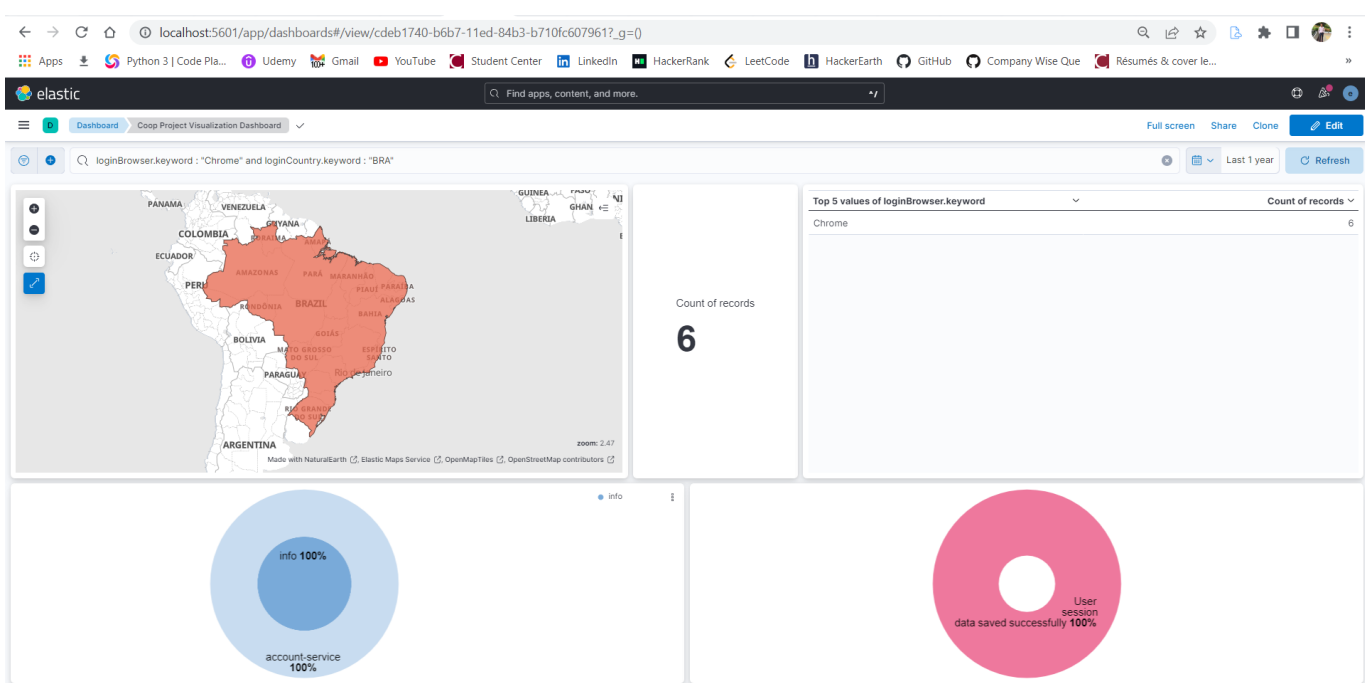
Finally, Kibana Query Language is a sophisticated and adaptable language that enables users to search, filter, and analyse data.

Some of the examples of the KQL on our Kibana Dashboard are as follows:

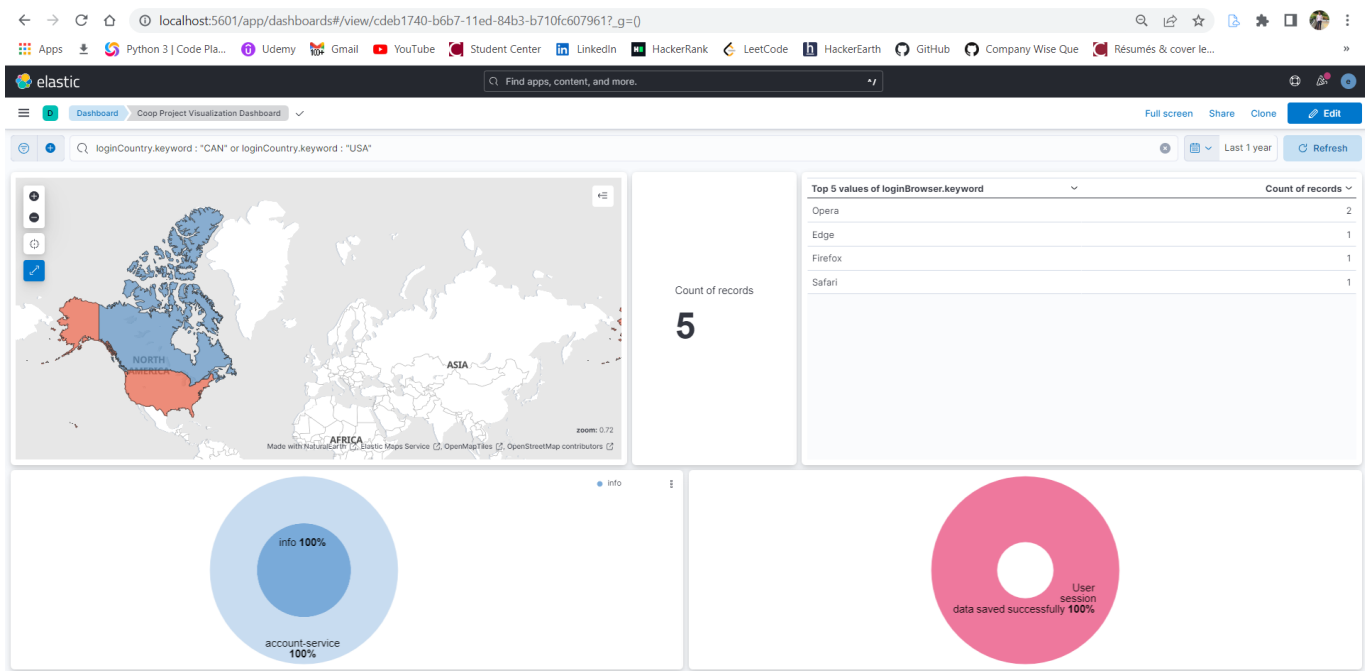
loginCountry.keyword : "CAN"



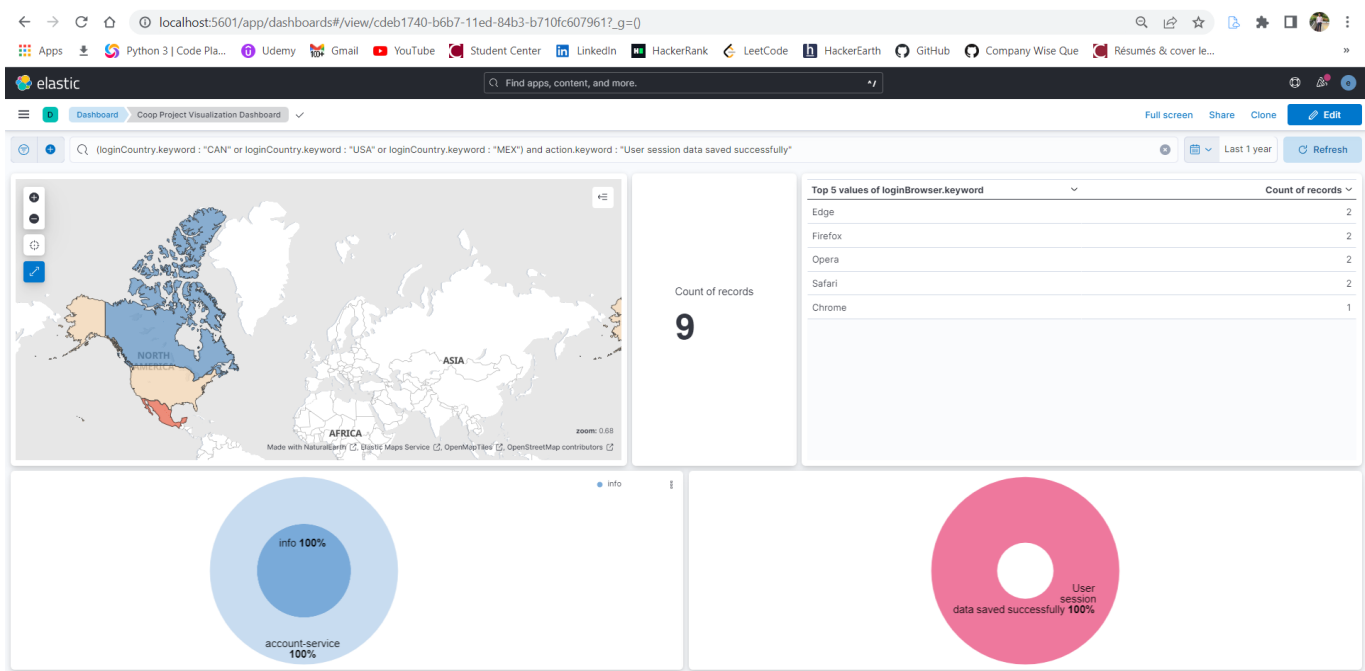
loginBrowser.keyword : "Chrome" and loginCountry.keyword : "BRA"



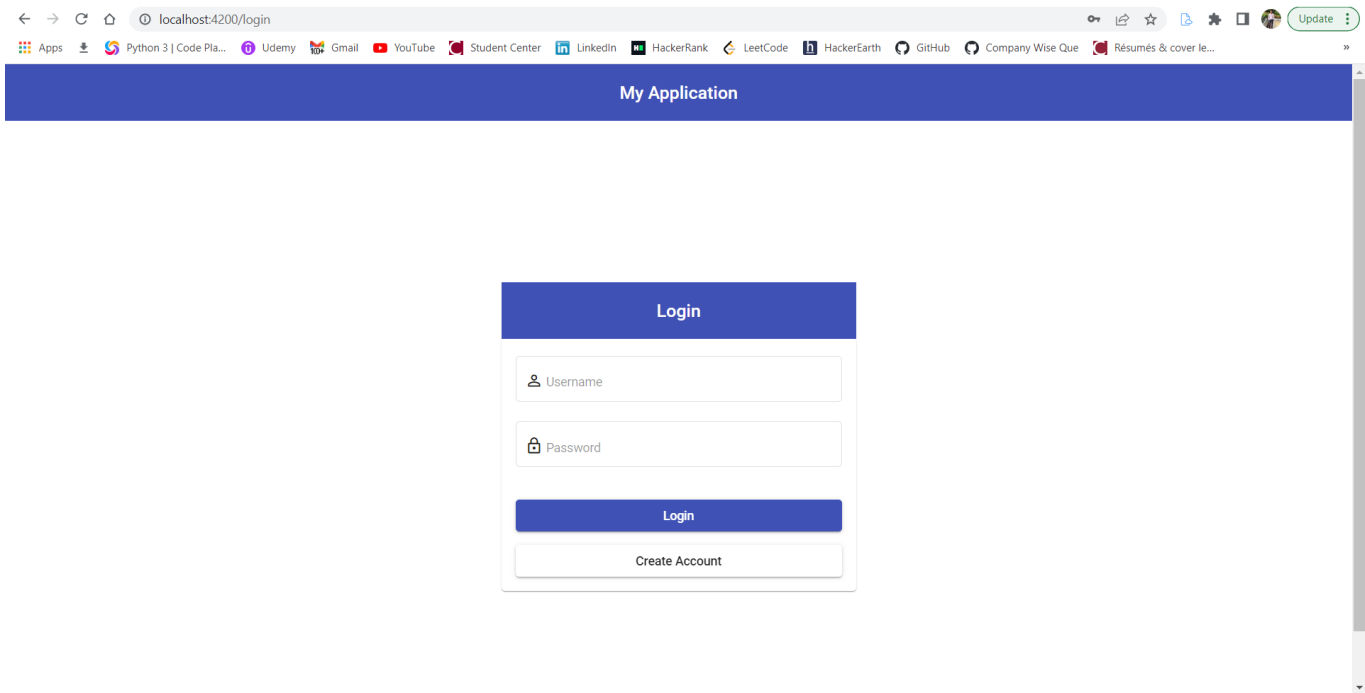
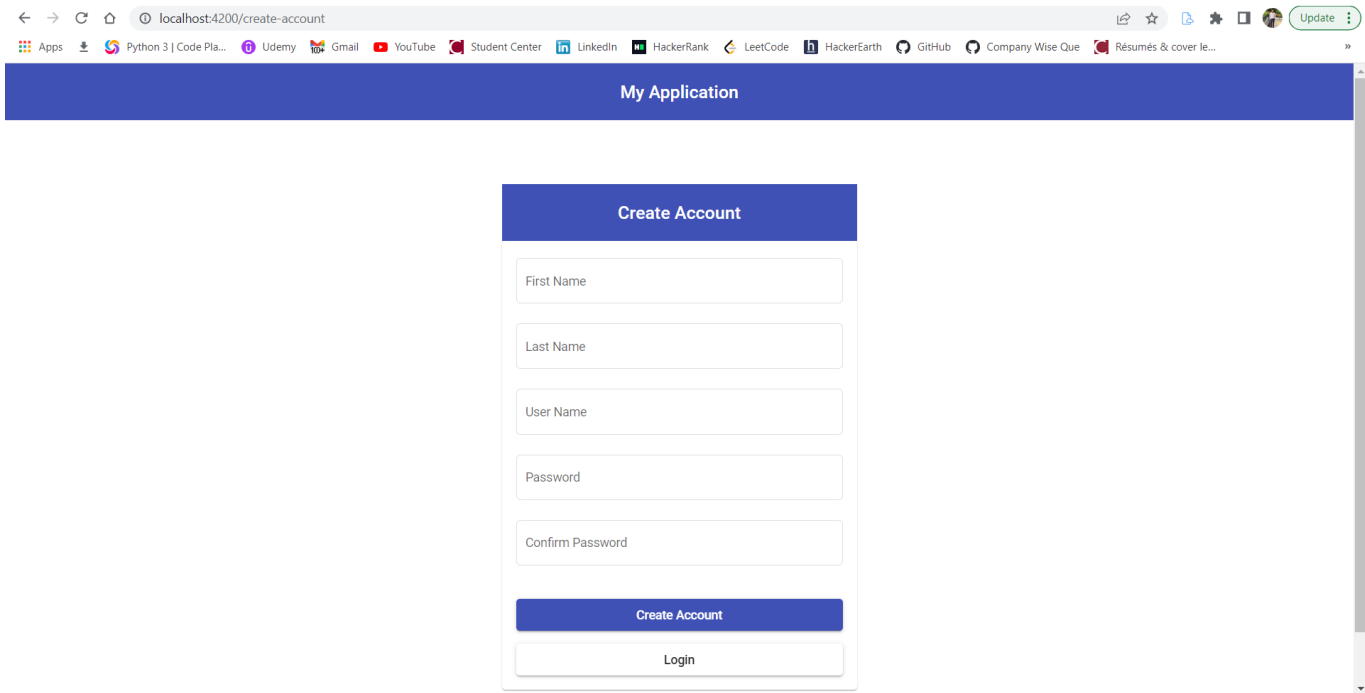
loginCountry.keyword : "CAN" or loginCountry.keyword : "USA"



(loginCountry.keyword : "CAN" or loginCountry.keyword : "USA" or loginCountry.keyword : "MEX") and action.keyword : "User session data saved successfully"



Screenshots of the Application



My Application

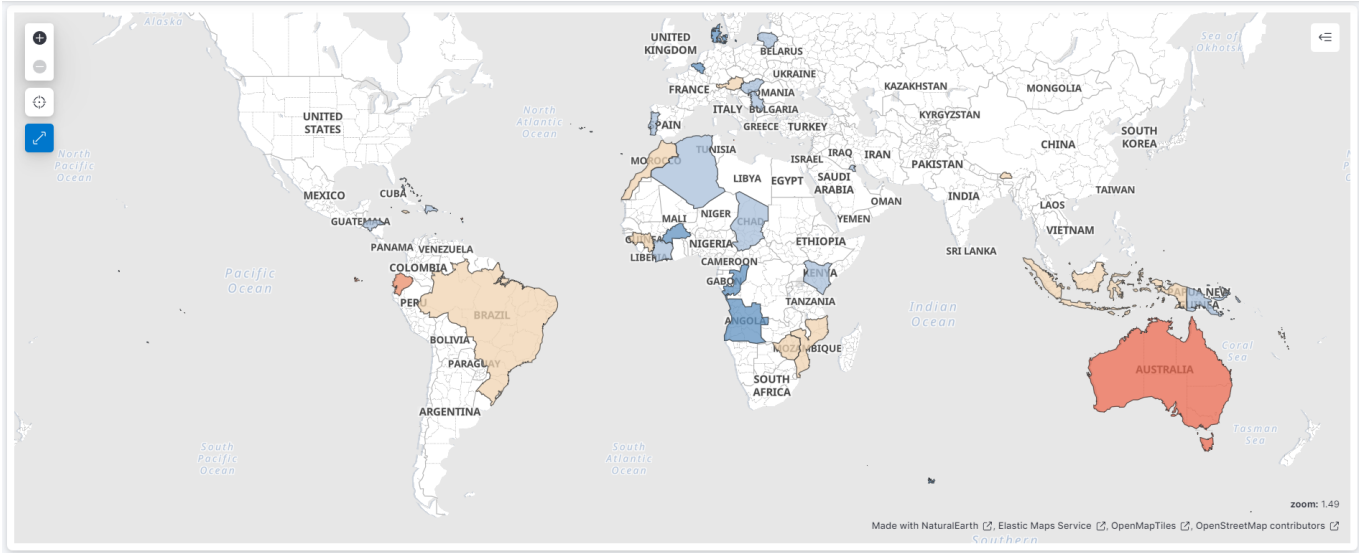
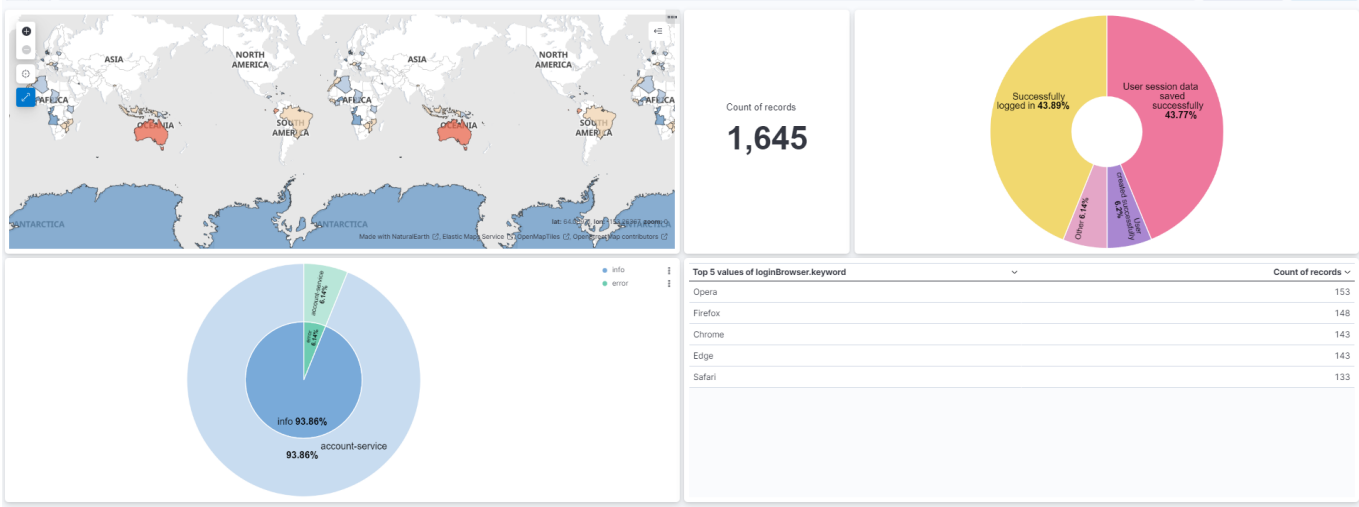
Dashboard

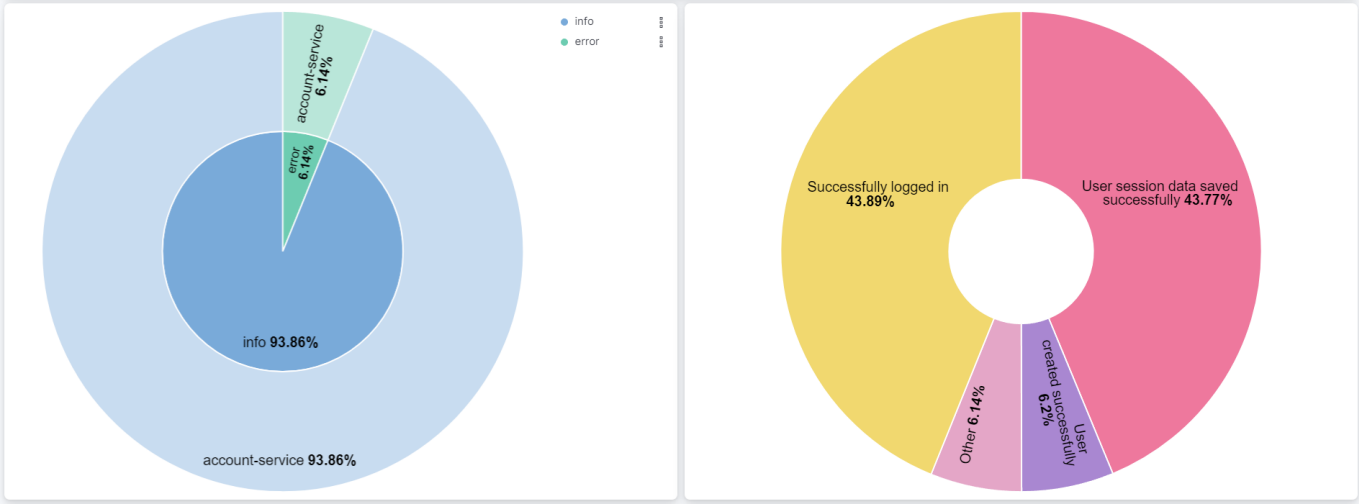
Your Current Session Start Time:

 Date: 25/21/2023 & Time: 18:30:43

Log Out

Screenshots of the Kibana Dashboard





Comparison with other existing Solutions

1) Comparison with Logtail

Logtail is a lightweight log shipping and management tool that focuses on simplicity and ease of use. It is designed to be a simple, reliable, and efficient tool for shipping logs from multiple sources to a central location for storage and analysis. Logtail can be easily configured to ship logs to popular cloud-based logging services like Datadog, Sumo Logic, and Loggly.

Logtail, owing to custom-built technology and ClickHouse-powered storage, provides a considerably more resource-efficient and hence cheaper and quicker alternative to the ELK stack. Logtail enables you to collect logs from throughout your stack, utilise the collected data to the greatest extent feasible, and so spend less time debugging and troubleshooting.

Utilizing the built-in Google Docs-like features. You may preserve key log pieces, comment on them, and share them with your colleagues. You may also use Logtail's presence and absence notifications to detect problems before they become serious. Logtail supports SQL-compatible queries and allows you to effectively query your logs or link them straight to your BI tool for bespoke reports.

Logtail also includes a live log tail, which allows you to look up logs in real time, scroll, and analyse data with a single click utilising its well-designed and purpose-driven dark mode Interface. Grafana handles all data visualisations, allowing you to construct bespoke and customised Charts and Dashboards.

Logtail can manage logs from any platform, and installation takes only a few minutes. Logtail's main advantages are collaboration features, ClickHouse-based storage, and the price beginning at \$0.25/GB.

Here are some key differences between ELK Stack and Logtail:

1. Complexity: ELK Stack is a more complex tool than Logtail, as it involves several components and requires more configuration and management. Logtail, on the other hand, is designed to be simple and easy to use, with minimal configuration required.
2. Scalability: ELK Stack is known for its ability to scale easily and handle large amounts of data. Logtail, on the other hand, may not be as scalable as ELK Stack but is still capable of handling moderate amounts of data.
3. Visualization: ELK Stack provides powerful visualization tools with Kibana, which allows users to explore and analyze log data in real time. Logtail, on the other hand, does not offer built-in visualization tools, but it can be easily integrated with other third-party visualization tools.
4. Cost: ELK Stack is an open-source tool and is therefore free to use, but it requires more resources and maintenance. Logtail, on the other hand, is a commercial product that offers a free trial but requires a subscription to use in production.

Overall, ELK Stack is a more powerful and feature-rich tool than Logtail, but it requires more resources and expertise to set up and maintain. Logtail is a simpler and more user-friendly tool

that is ideal for smaller-scale logging needs or for those who want a quick and easy way to ship logs to a central location.

2) Comparison with Sumo Logic

Sumo Logic is a cloud-based log management and analytics platform that offers real-time insights into log data. Sumo Logic provides a unified platform for collecting, analyzing, and visualizing log data from different sources, including cloud-based applications and infrastructure, as well as on-premises systems. Sumo Logic is known for its ease of use, scalability, and advanced analytics capabilities.

Sumo Logic provides a comprehensive collection of log management solutions for the full stack, whether in the cloud, on-premises, or hybrid. Consolidated data visualisation enables you to detect emerging patterns and disarm any mistakes before they occur or during damage control, allowing you to uncover the core cause more quickly.

Anomaly detection, outlier identification, and predictive analytics provide deep and comprehensive insights into the performance of your infrastructure. Sumo logic provides real-time visibility into cloud applications and infrastructure on Amazon, Azure, and GCP. In addition, you will have access to over 150 applications and native connectors for complete out-of-the-box insight into third-party technologies.

Sumo Logic delivers two dashboards: a live dashboard and an interactive dashboard. The live dashboard displays a plethora of real-time data in chronological sequence. It does not, however, offer the opportunity to see earlier data. The interactive dashboard comes into play here. The interactive dashboard allows you to get a comprehensive overview of events and trends, zoom in on graphs, and spot uncommon events. To focus on certain problems and exceptions in the future, you may filter for them. Sumo Logic's main advantages are its two dashboard modes, security monitoring, and threat detection features.

Here are some key differences between ELK Stack and Sumo Logic:

1. Complexity: ELK Stack can be complex to set up and maintain, as it requires the installation and configuration of multiple components. Sumo Logic, on the other hand, is a cloud-based service that is easy to set up and manage, with minimal configuration required.
2. Scalability: Both ELK Stack and Sumo Logic are highly scalable, but Sumo Logic is a cloud-based platform that can handle very large volumes of data without requiring additional infrastructure.
3. Customizability: ELK Stack is highly customizable and can be tailored to meet specific needs. Sumo Logic, on the other hand, is a managed service that offers a range of features and capabilities, but it may not be as customizable as ELK Stack.
4. Cost: ELK Stack is an open-source tool and is therefore free to use, but it requires more resources and maintenance. Sumo Logic, on the other hand, is a commercial product that charges based on the amount of data processed, with different pricing plans available.

Overall, ELK Stack and Sumo Logic are both powerful tools for log management and analysis, but they have different strengths and weaknesses. ELK Stack is highly customizable and can be tailored to meet specific needs, but it requires more resources and expertise to set up and maintain. Sumo Logic is a cloud-based service that is easy to use and highly scalable, but it may not be as customizable as ELK Stack. The choice between the two ultimately depends on the specific needs and priorities of the organization.

3) Comparison with Logic Monitor

LogicMonitor, is a cloud-based IT infrastructure monitoring and intelligence platform. It is designed to be a simple, automated, and intelligent way to monitor IT infrastructure, including servers, networks, applications, and cloud services. LogicMonitor provides real-time monitoring, alerting, and reporting tools to help users proactively identify and resolve issues.

For hybrid and multi-cloud settings, LogicMonitor provides log intelligence at scale. With an emphasis on data hygiene and corporate compliance, your data is consolidated, correlated, and contextualised. LogicMonitor enables you to consolidate your monitoring, correlate pertinent logs with metrics, and manage them all from a single platform.

It offers over 2000 on-premises and cloud connectors, modules, and pre-built templates. LogicMonitor is genuinely user-friendly since it provides query choices for users of various skill levels. It also gives you access to raw data that is up to 12 months old. Metrics, logs, and log anomalies are all related with the devices, cloud instances, and containers to which they are assigned.

LogicMonitor uses machine learning methods to alter your data, which reduces troubleshooting times and enables for improved productivity by relieving your engineers of unnecessary chores. Anomalies are discovered and contextualised automatically for simpler root cause investigation. LogicMonitor supports the whole IT operations lifecycle with connectors such as ServiceNow, CMDB, and Ansible.

One of the most significant drawbacks is the requirement to communicate your membership with a sales staff. You must obtain a personalised quotation. The primary advantages of LogicMonitor are as follows: Strong use of automation and machine learning approaches, suitable for all skill levels without sacrificing functionality

Here are some key differences between ELK Stack and LogicMonitor:

1. **Monitoring:** ELK Stack is primarily a log management tool, while LogicMonitor is a comprehensive IT infrastructure monitoring tool. LogicMonitor provides real-time monitoring and alerting for a wide range of IT infrastructure components, while ELK Stack is focused on log analysis and visualization.
2. **Deployment:** ELK Stack can be deployed on-premises or in the cloud, while LogicMonitor is a cloud-based tool that is accessed via a web browser. This means that LogicMonitor is easier to set up and use, while ELK Stack gives users more control over their infrastructure.
3. **Scalability:** Both ELK Stack and LogicMonitor are scalable tools, but LogicMonitor has the advantage of being a cloud-based tool, which allows it to easily scale up or down

depending on the user's needs. ELK Stack, on the other hand, requires more resources and expertise to scale effectively.

4. **Cost:** ELK Stack is an open-source tool and is therefore free to use, but it requires more resources and maintenance. LogicMonitor, on the other hand, is a commercial product that requires a subscription to use, but it includes maintenance, updates, and technical support.

Overall, ELK Stack is a powerful tool for log analysis and visualization, but it is not a comprehensive IT infrastructure monitoring tool like LogicMonitor. LogicMonitor provides real-time monitoring, alerting, and reporting tools for a wide range of IT infrastructure components, making it a more comprehensive tool for IT infrastructure monitoring. However, LogicMonitor is a commercial product, while ELK Stack is an open-source tool that requires more expertise and resources to set up and maintain.

4) Comparison with Loggly

Loggly is a cloud-based log management and analysis tool. It is designed to be a simple, scalable, and secure way to collect, store, and analyze log data. Loggly provides real-time analytics, alerts, and dashboards to help users gain insights into their log data.

SolarWinds' Loggly is a log management and aggregation solution. It is presently one of the most popular market solutions. Loggly is an agentless log analyzer that collects data from application servers directly. Loggly may retrieve data from pre-existing software by using a token or the standard Syslog using HTTP(s).

It supports different languages and systems and can deal with txt-based logs from any source. Ruby, Java, Python JavaScript, PHP, Apache HTTP Server, Tomcat, MySQL, Syslog-ng, rsyslog, and many other languages are supported. Loggly's major focus is on identifying and resolving operational issues. Loggly is a robust log analysis tool thanks to its customizable dashboards, documentation, and a wide range of useful functions. Loggly's Major Advantages: SolarWinds Support

Here are some key differences between ELK Stack and Loggly:

1. **Deployment:** ELK Stack can be deployed on-premises or in the cloud, while Loggly is a cloud-based tool that is accessed via a web browser. This means that ELK Stack gives users more control over their infrastructure, while Loggly is easier to set up and use.
2. **Scalability:** Both ELK Stack and Loggly are scalable tools, but Loggly has the advantage of being a cloud-based tool, which allows it to easily scale up or down depending on the user's needs. ELK Stack, on the other hand, requires more resources and expertise to scale effectively.
3. **Cost:** ELK Stack is an open-source tool and is therefore free to use, but it requires more resources and maintenance. Loggly, on the other hand, is a commercial product that requires a subscription to use, but it includes maintenance, updates, and technical support.
4. **Features:** Loggly includes several features that ELK Stack does not have out-of-the-box, such as dynamic field explorer, tag cloud visualizations, and live tailing. Loggly also provides more pre-built integrations with popular tools like Slack, PagerDuty, and JIRA.

Overall, ELK Stack is a more powerful and feature-rich tool than Loggly, but it requires more resources and expertise to set up and maintain. Loggly is a simpler and more user-friendly tool that is ideal for smaller-scale logging needs or for those who want a quick and easy way to collect, store, and analyze log data. However, Loggly is a commercial product, while ELK Stack is an open-source tool that requires more expertise and resources to set up and maintain.

5) Comparison with Sematext

Sematext is a cloud-based log management and monitoring tool. It is designed to be a simple, scalable, and secure way to collect, store, and analyze log data, as well as monitor performance metrics, user experience, and uptime. Sematext provides real-time analytics, alerts, and dashboards to help users gain insights into their log data.

Sematext is a service for monitoring and logging. It supports centralised logging, allowing you to collect and store logs from any data source in a single location. Data can be collected from servers, apps, databases, containers, systems, and other sources. Sematext enables you to examine your logs in real time as they arrive in the cloud from various data sources.

It employs the ELK stack for data collection and transformation, searching, filtering, and analysing, and finally data management and visualisation, but with the benefit of a sophisticated and hosted solution. With real-time alerts on metrics and logs, you can troubleshoot more quickly. To speed up the process, log analysis and anomaly detection are employed. Email, PagerDuty, Slack, HipChat, BigPanda, OpsGenie, VictorOps, WebHooks, Nagios, Zapier, and more services can be integrated.

Sematext is hosted on Amazon, whose infrastructure adheres to stringent IT security best practises. Your logs are encrypted using HTTPS and sent using TLS/SLL channels. Furthermore, you may limit particular rights to certain members of your team to improve the integrity and security of your service.

The primary advantages of Sematext are as follows: It combines infrastructure and application performance monitoring, as well as log management. Simple to use with well-configured dashboards and reports, making it easy to get started, and no need for extensive setting.

Here are some key differences between ELK Stack and Sematext:

1. **Deployment:** ELK Stack can be deployed on-premises or in the cloud, while Sematext is a cloud-based tool that is accessed via a web browser. This means that ELK Stack gives users more control over their infrastructure, while Sematext is easier to set up and use.
2. **Scalability:** Both ELK Stack and Sematext are scalable tools, but Sematext has the advantage of being a cloud-based tool, which allows it to easily scale up or down depending on the user's needs. ELK Stack, on the other hand, requires more resources and expertise to scale effectively.
3. **Cost:** ELK Stack is an open-source tool and is therefore free to use, but it requires more resources and maintenance. Sematext, on the other hand, is a commercial product that requires a subscription to use, but it includes maintenance, updates, and technical support.

4. Features: Sematext includes several features that ELK Stack does not have out-of-the-box, such as transaction tracing, anomaly detection, and distributed tracing. Sematext also provides more pre-built integrations with popular tools like Slack, PagerDuty, and Datadog.

Overall, ELK Stack is a more powerful and feature-rich tool than Sematext, but it requires more resources and expertise to set up and maintain. Sematext is a simpler and more user-friendly tool that is ideal for smaller-scale logging needs or for those who want a quick and easy way to collect, store, and analyze log data. However, Sematext is a commercial product, while ELK Stack is an open-source tool that requires more expertise and resources to set up and maintain.

6) Comparison with LogDNA

LogDNA is a cloud-based log management tool that is designed to be a simple, scalable, and secure way to collect, store, and analyze log data. LogDNA provides real-time analytics, alerts, and dashboards to help users gain insights into their log data.

LogDNA automatically parses main log line types and provides Custom Parsing Templates. You may filter your logs by app, host, or cluster, rapidly read logs from any source, and search through them using basic keywords, exclusion terms, chained expressions, and data ranges. Alerts may be set off based on presence or absence, or they can be generated from a stored View and reported on in PagerDuty, Slack, or using a custom Webhook. LogDNA also allows you to store and share views for quick access to popular Filters and Searches.

Elasticsearch is the foundation of LogDNA. Filtering, log grouping by source, and other tasks are handled through a web-based Interface. You may also interact with user-specific logs and use visualisation and custom dashboards. Agentless log collecting through Syslog and HTTP(s) is possible, including full-text search and visualisations.

LogDNA's pricing packages are based on the number of users and the retention term in days. To begin, LogDNA is available for free for one user with no log retention and unlimited stored views. LogDNA's main advantages are its pay-as-you-go pricing mechanism and well-designed user interface.

Here are some key differences between ELK Stack and LogDNA:

1. Deployment: ELK Stack can be deployed on-premises or in the cloud, while LogDNA is a cloud-based tool that is accessed via a web browser. This means that ELK Stack gives users more control over their infrastructure, while LogDNA is easier to set up and use.
2. Scalability: Both ELK Stack and LogDNA are scalable tools, but LogDNA has the advantage of being a cloud-based tool, which allows it to easily scale up or down depending on the user's needs. ELK Stack, on the other hand, requires more resources and expertise to scale effectively.
3. Cost: ELK Stack is an open-source tool and is therefore free to use, but it requires more resources and maintenance. LogDNA, on the other hand, is a commercial product that requires a subscription to use, but it includes maintenance, updates, and technical support.

4. Features: LogDNA includes several features that ELK Stack does not have out-of-the-box, such as Live Tail, anomaly detection, and integrations with popular tools like Slack, PagerDuty, and JIRA. ELK Stack, on the other hand, offers more advanced search and analysis capabilities, and its visualization tool, Kibana, is more powerful than LogDNA's built-in dashboard.

Overall, ELK Stack is a more powerful and feature-rich tool than LogDNA, but it requires more resources and expertise to set up and maintain. LogDNA is a simpler and more user-friendly tool that is ideal for smaller-scale logging needs or for those who want a quick and easy way to collect, store, and analyze log data. However, LogDNA is a commercial product, while ELK Stack is an open-source tool that requires more expertise and resources to set up and maintain.

7) Comparison with Dynatrace

Dynatrace is a commercial APM (Application Performance Management) tool that provides full-stack monitoring and observability for cloud-native applications. Dynatrace uses AI-powered automation to provide real-time insights into application performance, user experience, and infrastructure health. Dynatrace supports a wide range of technologies and platforms, including containers, microservices, cloud platforms, and more.

Log Monitoring from the Dynatrace portfolio gives you access to and monitoring of logs for all of your mission-critical activities. It is simple to create custom log metrics, which will allow you to monitor and interpret log data in the context of the rest of your infrastructure in real-time.

Logs may be filtered based on keywords or chronology and processed using AI, which links log entries with issues and leverages this association in root-cause analysis. If you utilise Dynatrace as a SaaS, you can choose Log Monitoring v1 or Log Monitoring v2. The documentation provided by Dynatrace covers all of the subtleties. Dynatrace, on the other hand, is more difficult to master and needs more time.

Dynatrace provides a full-stack monitoring solution as well as several individual plans. The AI-assisted full-stack monitoring solution, more than 560 supported technologies, and solutions that include security, digital experience, and even business analytics are the main advantages of Dynatrace.

Here are some key differences between ELK Stack and Dynatrace:

1. Scope: ELK Stack is primarily a log management and analysis tool, while Dynatrace is a comprehensive APM tool that provides full-stack monitoring and observability for cloud-native applications.
2. Deployment: ELK Stack can be deployed on-premises or in the cloud, while Dynatrace is a cloud-based tool that is accessed via a web browser. Dynatrace also offers on-premises deployment options, but they require more expertise and resources to set up and maintain.
3. Scalability: Both ELK Stack and Dynatrace are scalable tools, but Dynatrace has the advantage of being a cloud-based tool, which allows it to easily scale up or down depending on the user's needs. Dynatrace also uses AI-powered automation to provide more efficient and accurate monitoring and alerting.

4. Cost: ELK Stack is an open-source tool and is therefore free to use, but it requires more resources and maintenance. Dynatrace, on the other hand, is a commercial product that requires a subscription to use, but it includes maintenance, updates, and technical support.
5. Features: Dynatrace includes several features that ELK Stack does not have out-of-the-box, such as AI-powered root cause analysis, real user monitoring, and deep database monitoring. Dynatrace also provides more detailed insights into application performance and infrastructure health, making it easier to identify and fix issues.

Overall, ELK Stack is a powerful and flexible tool for log management and analysis, but it may require more expertise and resources to set up and maintain. Dynatrace, on the other hand, provides comprehensive APM capabilities with AI-powered automation, making it easier to monitor and troubleshoot cloud-native applications. However, Dynatrace is a commercial product, while ELK Stack is an open-source tool that is free to use.

8) Comparison with Graylog

Graylog is a free and open-source log management tool that provides centralized log management, real-time analysis, and alerting. Graylog consists of several components, including Graylog Server, Elasticsearch, and MongoDB. Graylog Server collects, processes, and stores logs, Elasticsearch is used for indexing and searching, and MongoDB is used for metadata storage. Graylog also includes a web interface for visualizing and analyzing data.

Graylog operates on several models. Graylog open - their open-source solution - Graylog Small Business, or Graylog Enterprise are your options. Graylog cloud, the last option, provides the same experience as Graylog Enterprise but is hosted on the cloud, saving you money on your own infrastructure.

Graylog is capable of processing logs from any data source, data display, and analysis. It is built on Elasticsearch and MongoDB. The Dashboard is made up of widgets, each of which provides you with distinct information derived from various data kinds. You can see counts, charts, graphs, views, and more.

Graylog's many deployment options allow you to operate and maintain it on your own or have it hosted, giving you greater freedom and control. The user interface is unquestionably more appealing. Graylog's websites, on the other hand, are hardly a designer's dream.

Graylog's main advantages are as follows: Even the free edition provides a wide range of functionality. The ability to search for multiple criteria without manually filtering the data, and There is an open-source solution available.

Here are some key differences between ELK Stack and Graylog:

1. Deployment: ELK Stack can be deployed on-premises or in the cloud, while Graylog can also be deployed on-premises or in the cloud. Both tools offer installation packages and container images to make deployment easier.

2. **Scalability:** Both ELK Stack and Graylog are scalable tools, but ELK Stack has an advantage in terms of scalability due to its distributed architecture and ability to handle large amounts of data.
3. **Ease of Use:** Graylog is known for its user-friendly interface and ease of use, making it easier for users to set up and configure. ELK Stack, on the other hand, may require more technical expertise to set up and configure, although it has many resources available for help and support.
4. **Features:** Graylog and ELK Stack have similar features, such as real-time analysis and alerting, but Graylog also includes features such as user management, content packs for easier configuration, and built-in support for some log types. ELK Stack has a wider range of plugins and integrations available.
5. **Cost:** Both Graylog and ELK Stack are open-source tools, which means they are free to use. However, there may be additional costs associated with deploying and maintaining these tools, such as hosting costs, storage costs, and support costs.

Overall, both ELK Stack and Graylog are powerful tools for log management and analysis, but they have different strengths and weaknesses. Graylog may be more user-friendly and easier to set up, while ELK Stack is known for its scalability and wider range of plugins and integrations. Ultimately, the choice between these tools will depend on the specific needs and requirements of the user.

9) Comparison with New Relic

New Relic is a cloud-based application performance monitoring (APM) tool that provides real-time performance monitoring, alerting, and analysis. New Relic offers a range of tools and features for monitoring applications, including transaction tracing, code profiling, and error tracking.

Infrastructure monitoring from New Relic allows for greater visibility and troubleshooting. With only a few clicks, New Relic provides an all-in-one data observation tool capable of correlation or drill-down from Kubernetes to specialised log tracing.

Because New Relic is extremely customizable, it makes no difference if you operate from one or more clouds on-premise; you will have access to particular, precise, and bespoke metrics in real-time and on an infinite scale. New Relic is an open and adaptable integration network that supports all of the most common integrations, including Amazon, Azure, GCP, MYSQL, NGINX, and Kafka. If you discover an unsupported integration, you may create one from scratch using NewRelic's Flex integration builder.

You may download New Relic for free and use its basic log management and analysis tools. The remaining options are charged based on your consumption, with you paying for everything you used more than the free plan. The main advantages of New Relic are Kubernetes monitoring Pixie and many solutions according on use scenario.

Here are some key differences between ELK Stack and New Relic:

1. **Deployment:** ELK Stack can be deployed on-premises or in the cloud, while New Relic is a cloud-based tool that is only available as a Software-as-a-Service (SaaS) solution.

2. Scope: ELK Stack is primarily focused on log management and analysis, while New Relic is primarily focused on application performance monitoring (APM).
3. Ease of Use: New Relic is known for its user-friendly interface and ease of use, making it easier for users to set up and configure. ELK Stack, on the other hand, may require more technical expertise to set up and configure, although it has many resources available for help and support.
4. Features: New Relic offers a range of features for monitoring applications, including transaction tracing, code profiling, and error tracking, while ELK Stack has a wider range of plugins and integrations available for log management and analysis.
5. Cost: New Relic is a paid tool, with pricing based on the number of hosts, applications, and services being monitored. ELK Stack is an open-source tool that is free to use, although there may be additional costs associated with deploying and maintaining the tool, such as hosting costs, storage costs, and support costs.

Overall, both ELK Stack and New Relic are powerful tools for monitoring and analyzing data, but they have different strengths and weaknesses. New Relic may be more user-friendly and easier to set up, while ELK Stack is known for its scalability and wider range of plugins and integrations. Ultimately, the choice between these tools will depend on the specific needs and requirements of the user, such as whether they are focused on log management or application performance monitoring, and whether they prefer a cloud-based or on-premises solution.

10) Comparison with Splunk

Splunk is a commercial log management and analysis tool that provides real-time data insights across applications, infrastructure, and security. Splunk offers a range of tools and features for searching, analyzing, and visualizing data, including machine learning and data integration capabilities.

Splunk is a fresh and cutting-edge log management and monitoring system. It is also available for mobile use and supports augmented reality.

Splunk also has functionality for searching, filtering, diagnosing, indexing, and reporting in addition to log management. It also has user-friendly dashboards that may be separated into various pertinent areas. Splunk employs distributed tracing, a technique for monitoring events, failures, and performance concerns.

Splunk searches for short-term data quickly. Nevertheless, it lags behind when collecting data over a longer period of time or detecting patterns. Splunk, on the other hand, offers a plethora of extra functionalities. Live logging, S3 backup, Heroku support, Github integration, JIRA integration, and more features are available.

Splunk's main advantages include support for numerous technologies such as S3 backup, live logging, Heroku, Github, and others, as well as a flexible Interface, query language support, and a complex, enterprise-ready solution.

ELK Stack is often compared to Splunk, which is a proprietary log management and data analysis tool. While both solutions offer similar capabilities, the ELK Stack is open-source and offers greater flexibility and customization. However, Splunk is known for its ease of use and intuitive interface and offers a wide range of pre-built visualizations and alerts. Both solutions have their own strengths and weaknesses, and the choice between them will depend on the specific needs and resources of the organization.

Here are some key differences between ELK Stack and Splunk:

1. **Deployment:** ELK Stack can be deployed on-premises or in the cloud, while Splunk can be deployed on-premises or as a cloud-based Software-as-a-Service (SaaS) solution.
2. **Scalability:** ELK Stack is known for its scalability, as it can easily handle large amounts of data and scale horizontally as needed. Splunk also has scalability capabilities, but may require additional infrastructure and licensing costs to handle large volumes of data.
3. **Ease of Use:** Splunk is known for its user-friendly interface and ease of use, making it easier for users to set up and configure. ELK Stack, on the other hand, may require more technical expertise to set up and configure, although it has many resources available for help and support.
4. **Features:** Splunk offers a wide range of features for log management and analysis, including machine learning and data integration capabilities, while ELK Stack has a wider range of plugins and integrations available for log management and analysis.
5. **Cost:** Splunk is a commercial tool with a licensing model based on the amount of data being ingested and indexed, which can be costly for large-scale deployments. ELK Stack is an open-source tool that is free to use, although there may be additional costs associated with deploying and maintaining the tool, such as hosting costs, storage costs, and support costs.

Overall, both ELK Stack and Splunk are powerful tools for log management and analysis, but they have different strengths and weaknesses. Splunk may be more user-friendly and offer more advanced features, while ELK Stack is known for its scalability and wider range of plugins and integrations. Ultimately, the choice between these tools will depend on the specific needs and requirements of the user, such as whether they prefer a commercial or open-source solution, and whether they are focused on log management or other data insights.

Conclusion

Coop point of view :

A Proof Concept is made on ELK Stack which will help the company to migrate from Splunk to ELK Stack. It will be above and beyond the scope of my internship.

Professional Growth point of view :

It helped to learn technologies which a Full-Stack developer is expected to know in order to get a full-time job and learn ELK stack which is an open-source technology, due to which many companies are moving to it from costly tools like Splunk, New Relic, etc and many job opening's descriptions state about it too. It also gave me a chance to showcase my this technical skills with a live demo to my internship company manager in order to give a positive & strong image which will help me to convert my coop to a full-time job opportunity on top of learning very important technical skills which are considered hot technologies in the market.

Overall, Implementing a project utilising the ELK (Elasticsearch, Logstash, and Kibana) stack for log debugging and analysis, in conjunction with a full-stack application, was a good opportunity for me to demonstrate my talents to my management. In doing so, I exhibited not just my technical expertise but also my enthusiasm to acquire new technologies and my capacity to go above and beyond the requirements of my coop job.

It also aided in the implementation of an open-source ELK stack solution for log analysis, which will save the organisation money because it is license-free and will aid in the transition from pricey programmes such as Splunk. Utilizing the ELK stack assisted me in centralising and analysing logs created by various apps and servers, providing me with vital insights into the functioning of my system and spotting any issues that developed. I was able to demonstrate my abilities to design, create, and deploy a complete system, from the front-end to the back-end, by creating a full-stack application.

In conclusion, creating a project utilising the ELK stack and a full-stack application was an excellent method for me to exhibit my talents, initiative, and want to learn. As a result, I enhanced my visibility to my boss, boosted my prospects of converting my coop position to full-time employment, and positioned myself for future career advancement opportunities.

References

1. <https://betterstack.com/community/comparisons/elk-stack-alternatives/>
2. <https://signoz.io/blog/elk-alternatives/>
3. <https://www.gartner.com/reviews/market/application-performance-monitoring-and-observability/vendor/elasticsearch/product/elastic-elk-stack/alternatives?marketSeoName=application-performance-monitoring&vendorSeoName=elasticsearch&productSeoName=elastic-elk-stack>
4. <https://seoaves.com/datadog-competitors-and-alternatives-analysis/>
5. <https://www.rapid7.com/blog/post/2014/04/04/logs-as-data-what-are-the-top-use-cases-for-your-logs/>
6. <https://en.wikipedia.org/wiki/MongoDB>
7. <https://www.mongodb.com/>
8. <https://en.wikipedia.org/wiki/Express.js>
9. <https://expressjs.com/>
10. [https://en.wikipedia.org/wiki/Angular_\(web_framework\)](https://en.wikipedia.org/wiki/Angular_(web_framework))
11. <https://angular.io/>
12. <https://en.wikipedia.org/wiki/Node.js>
13. <https://nodejs.org/>
14. [https://en.wikipedia.org/wiki/Cypress_\(software\)](https://en.wikipedia.org/wiki/Cypress_(software))
15. <https://www.cypress.io/>
16. <https://github.com/winstonjs/winston>
17. <https://betterstack.com/community/guides/logging/how-to-install-setup-and-use-winston-and-morgan-to-log-node-js-applications/>
18. <https://en.wikipedia.org/wiki/Elasticsearch>
19. <https://www.elastic.co/>
20. <https://logz.io/blog/logstash-tutorial/>
21. <https://www.elastic.co/logstash>
22. <https://github.com/elastic/logstash>
23. <https://en.wikipedia.org/wiki/Kibana>
24. <https://www.elastic.co/kibana>
25. <https://github.com/elastic/kibana>
26. <https://github.com/elastic/elasticsearch>
27. <https://www.splunk.com/>
28. <https://newrelic.com/>
29. <https://www.graylog.org/>
30. <https://www.dynatrace.com/>
31. <https://www.mezmo.com/>
32. <https://sematext.com/>
33. <https://www.loggly.com/>
34. <https://www.logicmonitor.com/>
35. <https://www.sumologic.com/>
36. <https://betterstack.com/logtail>
37. <https://logz.io/learn/complete-guide-elk-stack/#intro>
38. <https://www.elastic.co/guide/en/elasticsearch/reference/current/tune-for-indexing-speed.html>
39. <https://logz.io/blog/kibana-tutorial/>

40. <https://www.learnsplunk.com/splunk-pricing---splunklicensing-model.html>
41. <https://www.crowdstrike.com/cybersecurity-101/observability/centralized-logging/>
42. <https://medium.com/@maheshd7878/elkfor-centralise-logging-d72aeaf8480>
43. <https://www.guru99.com/elk-stack-tutorial.html#2>
44. <https://www.elastic.co/what-is/elk-stack>
45. <https://www.digitalocean.com/community/tutorials/how-to-install-elasticsearch-logstash-and-kibana-elk-stack-on-ubuntu-14-04>
46. <https://www.elastic.co/what-is/kibana>
47. <https://www.elastic.co/guide/en/kibana/master/kueryquery.html#kuery-query>
48. <https://www.elastic.co/blog/using-parallel-logstashpipelines-to-improve-persistent-queue-performance>
49. <https://www.adservio.fr/post/scalability-patterns-forelk>
50. <https://www.elastic.co/guide/en/kibana/current/kuery-query.html>